

Agree to Disagree? A Tiered Taxonomy of Disagreement for Pluralistic Alignment

Anonymous ACL submission

Abstract

A substantial body of work applies deep learning to modeling human morality, treating rater agreement as evidence of a shared moral intuition, yet this agreement reflects how benchmarks filter out ambiguous and contested cases as much as it reflects any underlying moral substrate. Even work that treats persistent disagreement as a permanent feature of diverse societies, rather than a defect awaiting convergence, stops short of asking whether such diversity shares one underlying structure at all. We argue it does not: using a joint-holdability test, whether two positions can both be held without contradiction once fully understood, we introduce a tiered taxonomy distinguishing disagreement that resolves with shared context, disagreement that persists but remains representable side by side, and disagreement that resists comparison altogether, and show that alignment tools correctly aimed at each tier fail by quietly reintroducing the comparison they were built to avoid. We close without a fix, asking whether alignment work can learn to notice when a tool’s claimed target and its actual target diverge.

1 Introduction

A substantial body of work now applies deep learning to modeling human morality. Benchmarks such as ETHICS and Jiminy Cricket evaluate systems against human judgments, treating rater agreement as evidence that a moral intuition has been captured (Hendrycks et al., 2021a,b). But ambiguous and contested scenarios are typically filtered out during dataset construction, with inter-annotator agreement used as an inclusion criterion rather than an outcome to be explained; what these systems learn may be the subset of moral judgment already engineered for low variance, not a universal substrate awaiting discovery.

This matters beyond morality’s laboratory cases. Recent work argues that treating persistent dis-

agreement as a temporary defect rests on an Axiom of Rational Convergence that is optional and largely unsupported (Leibo et al., 2025), proposing instead that disagreement is the normal condition of a diverse society. This makes the comparability current alignment methods quietly assume, whether as a scalar reward or a written constitution, harder to defend. Pluralistic alignment work shares this premise and offers strategies such as Overton and steerable pluralism for representing competing values once diversity is acknowledged (Sorensen et al., 2024). Our question is prior to theirs: whether such diversity shares one underlying structure that a single representational strategy could handle.

This paper argues comparability fails, but not uniformly: some disagreement dissolves with shared context; some survives understanding yet remains representable side by side, unranked; and some resists comparison altogether, for want of a common standard. We introduce a tiered taxonomy built on these three intuitions, and show that tools correctly aimed at a given tier tend to fail the same way: by quietly re-flattening the difference they were built to preserve.

2 Related Work

Comparability

Preferentism, the assumption implicit across Reinforcement Learning from Human Feedback (RLHF) and much of Expected Utility Theory that any two options can in principle be ranked against a common scale, is the quiet premise beneath most current alignment methods. Zhi-Xuan et al. (2024) already challenge this view, arguing that reducing human values to preference satisfaction conflates the observable what of a choice with the why behind it. Edelman et al. (2025) build on this critique directly, proposing thick models of value, structured representations grounded in explicit normative standards tied to social role and context, as the

	Tier 1: Resolvable-Perspectival	Tier 2: Interpersonally Incomparable	Tier 3: Doctrinally Incommensurable
Analogy	6 to you, 9 to me	One man’s trash, another man’s treasure	Agree to disagree
Implications	Shared context is sufficient	Both valuations remain representable side by side	No representation resolves the disagreement
Correctly scoped tool	AI mediated deliberation	Pareto preserving multi-objective RL	Constitutional AI, disclosed as non-neutral
Within tier failure	Aggregation	Scalarization	Presenting a declared choice as neutral

Table 1: Summarized taxonomy of pluralistic disagreement (see Table 2 in Appendix).

alternative that preference orderings and unstructured text cannot support. This paper builds on that critique rather than departing from it. Where Zhi-Xuan et al.’s and Edelman et al.’s target is the reduction of value to preference, ours is the assumption, surviving even some non-preferentialist alternatives, that whatever underlies a choice is always comparable to whatever underlies another.

This claim is worth distinguishing from a related and better established one. Moral uncertainty (MacAskill et al., 2020) holds that we may not know which of several competing ethical theories is correct, an epistemic gap that further argument or evidence could in principle close. Our claim is structural rather than epistemic. For some comparisons, no privileged conversion exists between two positions even under complete information about both. Chang (2002) formalizes a version of this as parity, a fourth relation, alongside better, worse, and equal, in which two options are evaluable on a shared basis yet neither ranks above the other. Preferentialism has no place for parity. It assumes that wherever comparison is possible at all, ranking must eventually follow.

Morals

Work specifically aimed at modeling moral judgment offers an instructive case study here. Hendrycks et al. (2021a) introduce the ETHICS benchmark, which trains models against human judgments spanning several established ethical frameworks, including deontology, virtue ethics, and utilitarianism, filtering for scenarios where these frameworks converge. Hendrycks et al. (2021b) extend this approach to interactive text environments with the Jiminy Cricket benchmark, explicitly noting that their framework should not be read as representing absolute or unchanging moral judgment, and inviting future work to build frame-

works that might supersede it. Both projects are candid about this limitation, yet the benchmarks that result from it are, by construction, weighted toward the cases where the limitation matters least. Separately, work on RLHF has documented how readily the resulting systems drift from the judgments they are meant to encode, whether through reward hacking (Casper et al., 2023) or through sycophancy toward a rater’s apparent views (Perez et al., 2022). Constitutional AI (Bai et al., 2022) was proposed partly in response to these failures, replacing scalar feedback with a written set of principles, a move we return to directly in Section 5.

3 Beyond Morals

The moral cases discussed above serve mainly as an entry point rather than this paper’s real target: precedent for taking value plurality seriously, borrowed from a literature the field already accepts. The comparability assumption is easier to see, and harder to defend, once morality’s usual laboratory cases give way to domains where positions carry their own internal justification rather than being adjudicated by a shared theory of harm.

A separate line of work examines how **culture** is treated inside alignment pipelines. Recent research argues that current practice tends to reduce culture to static demographic categories or superficial facts, such as a country label or a short list of stated customs, rather than the deeper, often unstated meanings that give a practice its significance within a community (Orlowski et al., 2025). This is precisely the gap thick models of value are meant to close. Rather than representing values as scalar preferences or unstructured text, Edelman et al. (2025) propose structuring value representations so a system can distinguish enduring, socially embedded commitments from fleeting or superficial signals, and reason about them normatively

rather than merely recording them. A demographic label is exactly the kind of thin signal this approach is built to move past. **Religion** sits close to the center of this problem, since religious conviction is rarely reducible to a demographic label without losing the internal justification that makes it a source of genuine disagreement rather than mere preference. This is as one of the clearest patches in the quilt [Leibo et al. \(2025\)](#), in particular a piece of the pattern that cannot be pressed flat into a single converged fabric without erasing what gives it its shape.

Once named, this reduction is easy to notice elsewhere. **Social norms**, such as directness in conversation or expectations around hospitality, carry tacit justifications no demographic label can capture. **Lifestyle and taste** sit at a lower-stakes register of the same problem. **Politics** spans both registers, since political conviction often rests on the same kind of comprehensive worldview as religious belief. Together, these domains, considerably wider than the moral dilemmas that opened the discussion, are this paper's actual territory.

4 A Tiered Taxonomy of Disagreement

If the comparability assumption breaks down, it does not do so all at once. We distinguish three tiers, each named by a familiar intuition and each defined by what happens to the disagreement under conditions of full mutual understanding, a test we call joint holdability which assesses whether both original claims still be held without contradiction once both parties fully understand each other's position. Further case studies may be found in [Appendix A.1](#).

Tier 1: Resolvable Perspectival

Tier 1 disagreement is illustrated by the phrase "**six to you, nine to me**". Two observers, viewing the same numeral painted on the ground from opposite sides of it, report different readings, one seeing a six, the other a nine. Neither observer is wrong. Once they walk around and compare positions, both descriptions survive intact, folded into a single larger picture that neither had to abandon, a shape that reads six from one side and nine from the other. This is disagreement's null case. It looks like pluralism from either single vantage point, since each observer is reporting something true and the two reports conflict on their face, but it passes the joint holdability test as soon as the vantage points

are shared. What resolves the disagreement is information rather than persuasion. Neither party changes what they value or believe, only what they know about the other's position within the same larger picture.

The practical signature of a Tier 1 disagreement is that it **dissolves through explanation alone**. No negotiation is required, no side is asked to give anything up, and no representational device is needed to hold both original claims together once the missing context is supplied.

Tier 2: Interpersonally Incomparable

Tier 2 disagreement is illustrated by the phrase "**one man's trash is another man's treasure**". Two people assign opposing values to the same object, and no amount of mutual understanding produces a common scale to adjudicate between them. Consider a single household purchase, an expensive coat bought by one partner and viewed by the other. To the partner who bought it, the coat represents a well earned reward after a difficult year, an act of self expression rather than waste. To the other partner, the same coat represents money that should have gone toward savings or shared obligations, an act of indulgence rather than reward. Both partners can fully explain their reasoning to the other. Both accounts are internally coherent, and neither party is factually mistaken about anything, including the other's reasoning. What is missing is not information but a shared unit in which the value of self expression and the value of financial prudence could be weighed against each other.

This is [Robbins \(1932\)](#)'s classical objection to interpersonal utility comparison, applied here beyond economics and beyond a single household. There is no way to verify that one person's valuation and another's are even measured in commensurable units, so ranking them is not merely difficult but the wrong kind of operation to attempt. This is structurally the move [Arrow \(1951\)](#)'s impossibility theorem was built to avoid relying on. It is also, not incidentally, the operation a reward model performs every time it aggregates many annotators' pairwise judgments into a single scalar reward, treating disagreements of this kind as though they were noisy estimates of one true preference rather than two genuinely different, equally coherent valuations. Where a Tier 2 disagreement is correctly recognized, the honest response is not to search harder for the missing scale, but to **represent both valuations side by side**, as a distribution or a frontier

rather than a single winner.

Tier 3: Doctrinally Incommensurable

Tier 3 disagreement is illustrated by the phrase "agree to disagree". Two comprehensive world-views conflict not because a shared scale is missing, but because **no shared standard of adjudication exists** at all. Full mutual understanding does not soften this kind of disagreement, it clarifies exactly where it sits. This is close to what Rawls (1993) calls the burdens of judgment, the idea that reasonable people, reasoning in good faith under free institutions, can still arrive at irreconcilable comprehensive doctrines, not from bias or missing evidence, but because evidence is genuinely complex and a person's total life experience shapes how both evidence and value are weighed in ways further argument does not simply overwrite.

Religious conviction is the clearest case. Two people can each fully understand the other's theology, its internal logic, and its emotional weight, and still hold positions that cannot both be true. Neither party is asserting a claim that is true from where they stand. Both are asserting a claim they take to be simply true, which is exactly what a bigger shared picture, of the kind available in Tier 1, cannot contain.

5 Where the Fixes Break, Tier by Tier

While each tier admits a well motivated response, but examined closely each response fails in a way that quietly reintroduces the flattening it was designed to avoid.

Tier 1: Aggregation. The Habermas Machine (Tessler et al., 2024) was built to help groups reach consensus by generating group statements from participants' opinions and critiques, explicitly optimizing for predicted group approval. The system trains a personalized reward model on each participant's rankings, then selects among candidate statements using the Schulze method, a voting rule. Within its own scope, this is a legitimate Tier 1 tool. But the Schulze method is majoritarian: it can override a minority's preference to satisfy even a slight majority, a property its own designers acknowledge as a risk of minority-voice suppression. A statement that wins a vote is not the same as a statement every participant can hold without contradiction; the tool optimizes for approval, not for joint-holdability.

Tier 2: Scalarization. Multi-objective ap-

proaches that preserve a Pareto front rather than collapsing it to one scalar (Vamplew et al., 2008) are the correct Tier 2 response. But any system built this way must eventually act, and action selection typically requires a scalarization step, some rule for combining the multiple objectives back into a single choice. Any such rule has to fix, in advance, how much one objective is allowed to count against another, which is exactly the missing common scale a Tier 2 disagreement is defined by not having. Choosing that rule is not a technical detail to be engineered around, it is the same act of imposing comparability the tier was built to resist.

Tier 3: Disclosing a choice as neutral. Constitutional AI (Bai et al., 2022), understood as an openly disclosed, non-neutral reference frame, is a defensible Tier 3 move. Anthropic's own attempt to strengthen this disclosure, Collective Constitutional AI (Huang et al., 2024), sought public input through structured deliberation. The resulting high agreement and low polarization were reported as evidence of legitimacy, while the authors themselves note that the participant sample was small and not globally representative. The failure here is not that a choice was made, but disclosing the tool as a default, neutral choice rather than as one choice among possibilities; disclosure that an axis was chosen does not by itself establish that it is the right choice for those who never chose it.

6 Closing: Open Question, Not a Fix

Read together, these three failures share a shape. Each method built to preserve genuine plurality is eventually compressed into a single output: a winning statement, a selected point on a frontier, or one constitution presented as the system's voice, quietly forcing a false consensus onto disagreement that never resolved. Tiers 1 and 2 optimize a proxy until it no longer tracks its target: Goodhart (1975)'s Law beyond reward models. Tier 3 fits more loosely, with no optimization loop, only an imperfect proxy for legitimacy mistaken for legitimacy itself. We note the risk in naming all three this way: collapsing three mechanisms under one label is itself a small act of the flattening this paper describes. We do not propose a fix. What we leave open is narrower: whether alignment work can learn to notice when a tool's claimed target and its actual target have quietly come apart, and what it would mean to build systems that hold that difference rather than optimize it away.

References

Kenneth J. Arrow. 1951. *Social Choice and Individual Values*. John Wiley & Sons, New York.

Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, Andy Jones, Anna Chen, Anna Goldie, Azalia Mirhoseini, Cameron McKinnon, Carol Chen, Catherine Olsson, Christopher Olah, Danny Hernandez, Dawn Drain, Deep Ganguli, Dustin Li, Eli Tran-Johnson, Ethan Perez, and 32 others. 2022. [Constitutional ai: Harmlessness from ai feedback](#). *Preprint*, arXiv:2212.08073.

Stephen Casper, Xander Davies, Claudia Shi, Thomas Krendl Gilbert, J  r  my Scheurer, Javier Rando, Rachel Freedman, Tomasz Korbak, David Lindner, Pedro Freire, Tony Wang, Samuel Marks, Charbel-Rapha  l Segerie, Micah Carroll, Andi Peng, Phillip Christoffersen, Mehul Damani, Stewart Slocum, Usman Anwar, and 13 others. 2023. [Open problems and fundamental limitations of reinforcement learning from human feedback](#). *Preprint*, arXiv:2307.15217.

Ruth Chang. 2002. The possibility of parity. *Ethics*, 112(4):659–688.

Joe Edelman, Tan Zhi-Xuan, Ryan Lowe, Oliver Klingefjord, Vincent Wang-Mascianica, Matija Franklin, Ryan Othniel Kearns, Ellie Hain, Atrisha Sarkar, Michiel Bakker, Fazl Barez, David Duvenaud, Jakob Foerster, Iason Gabriel, Joseph Gubbels, Bryce Goodman, Andreas Haupt, Jobst Heitzig, Julian Jara-Ettinger, and 14 others. 2025. [Full-stack alignment: Co-aligning ai and institutions with thick models of value](#). *Preprint*, arXiv:2512.03399.

Charles A. E. Goodhart. 1975. Problems of monetary management: The UK experience. *Papers in Monetary Economics*, 1.

Dan Hendrycks, Collin Burns, Steven Basart, Andrew Critch, Jerry Li, Dawn Song, and Jacob Steinhardt. 2021a. Aligning AI with shared human values. In *International Conference on Learning Representations (ICLR)*.

Dan Hendrycks, Mantas Mazeika, Andy Zou, Sahil Patel, Christine Zhu, Jesus Navarro, Dawn Song, Bo Li, and Jacob Steinhardt. 2021b. What would Jiminy Cricket do? Towards agents that behave morally. In *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks (NeurIPS)*.

Saffron Huang, Divya Siddarth, Liane Lovitt, Thomas I. Liao, Esin Durmus, Alex Tamkin, and Deep Ganguli. 2024. [Collective constitutional AI: Aligning a language model with public input](#). In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency (FAccT ’24)*, pages 1395–1417, New York, NY, USA. ACM.

Joel Z. Leibo, Alexander Sasha Vezhnevets, William A. Cunningham, S  bastien Krier, Manfred Diaz, and

Simon Osindero. 2025. Societal and technological progress as sewing an ever-growing, ever-changing, patchy, and polychrome quilt. *arXiv preprint arXiv:2505.05197*.

William MacAskill, Krister Bykvist, and Toby Ord. 2020. *Moral Uncertainty*. Oxford University Press.

Eric J. W. Orlowski, Hakim Norhashim, and Tristan Koh Ly Wey. 2025. [’too much alignment; not enough culture’: Re-balancing cultural alignment practices in llms](#). *Preprint*, arXiv:2509.26167.

Ethan Perez, Sam Ringer, Kamil   Luko  i  t  , Karina Nguyen, Edwin Chen, Scott Heiner, Craig Pettit, Catherine Olsson, Sandipan Kundu, Saurav Kadavath, Andy Jones, Anna Chen, Ben Mann, Brian Israel, Bryan Seethor, Cameron McKinnon, Christopher Olah, Da Yan, Daniela Amodei, and 44 others. 2022. [Discovering language model behaviors with model-written evaluations](#). *Preprint*, arXiv:2212.09251.

John Rawls. 1993. *Political Liberalism*. Columbia University Press, New York.

Lionel Robbins. 1932. *An Essay on the Nature and Significance of Economic Science*. Macmillan, London.

Taylor Sorensen, Jared Moore, Jillian Fisher, Mitchell Gordon, Niloofar Mireshghallah, Christopher Michael Rytting, Andre Ye, Liwei Jiang, Ximing Lu, Nouha Dziri, Tim Althoff, and Yejin Choi. 2024. [A roadmap to pluralistic alignment](#). *Preprint*, arXiv:2402.05070.

Michael Henry Tessler, Michiel A. Bakker, Daniel Jarrett, Hannah Sheahan, Martin J. Chadwick, Raphael Koster, Georgina Evans, Lucy Campbell-Gillingham, Tantum Collins, David C. Parkes, Matthew Botvinick, and Christopher Summerfield. 2024. [AI can help humans find common ground in democratic deliberation](#). *Science*, 386(6719):eadq2852.

Peter Vamplew, John Yearwood, Richard Dazeley, and Adam Berry. 2008. [On the limitations of scalarisation for multi-objective reinforcement learning of Pareto fronts](#). In *AI 2008: Advances in Artificial Intelligence (Proceedings of the 21st Australasian Joint Conference on Artificial Intelligence)*, pages 372–378, Berlin, Heidelberg. Springer.

Tan Zhi-Xuan, Micah Carroll, Matija Franklin, and Hal Ashton. 2024. [Beyond preferences in AI alignment](#). *Philosophical Studies*. ArXiv:2408.16984.

A Appendix

A.1 Taxonomy Sample Case Studies

Disclaimer: The case studies below are illustrative rather than exhaustive or conclusive. They are meant to convey the flavor of how the taxonomy applies across different domains, not to offer a complete or authoritative classification of every disagreement a given tier or domain can produce. Several examples are also cross domain rather than cleanly belonging to one category, tipping, for instance, is framed here as a social norm, but is also shaped by local culture and, in places, by economic policy around service wages, and could reasonably be discussed under more than one heading. The domain label attached to each case study should be read as a convenient entry point rather than a firm boundary. The examples are offered in service of appreciating the taxonomy itself, how a disagreement's tier is determined by its logical structure, not by which domain it happens to be filed under, rather than as a settled taxonomy of domains in their own right.

Tipping Culture as Tier 1. Consider a diner from a country where tipping is not customary, since service staff are paid a full wage, eating at a restaurant in a country where tipping is expected, since staff wages are structured to be supplemented by it. The diner leaves no tip, intending no disrespect, simply following the convention they know. The server reads the missing tip as a judgment on the service, or as simple stinginess. Both readings feel obviously correct from where each party stands. But once the underlying convention is explained, the disagreement dissolves entirely: the diner was not being cheap, and the server was not being paranoid, each was reasoning correctly within a norm the other did not know existed. Nothing about what either party values needs to change, only what they know about how service work is compensated where the other is standing. This is Tier 1 in its purest form, the friction is real, but it is a gap in shared information, not a gap in shared values.

Aesthetics as Tier 2. Tier 2 disagreement is not confined to culture. Consider two people standing before the same abstract painting: one finds it genuinely moving, the color and form evoking something she cannot fully put into words but recognizes as significant; the other finds it unaccomplished, and values instead a technically demanding, representational work whose skill is immediately legible. Each can explain their standard fully to the other,

evocative ambiguity against demonstrated technical mastery, and neither standard is mistaken, nor is either viewer withholding information or reasoning carelessly. What is missing is not further explanation but a shared aesthetic currency in which the two standards could be weighed against each other to produce a single verdict. As with the household example, the honest response is not to keep arguing until one side concedes, but to hold both standards of value as legitimate and unranked.

Politics as Tier 3. Consider a disagreement over how much a state should redistribute individual wealth to fund shared social provision. One person reasons from a foundational commitment to individual liberty, holding that what a person earns is fundamentally theirs to keep, and that redistribution without consent, however well intentioned, is a form of coercion. The other reasons from a foundational commitment to collective welfare, holding that a person's earnings are only possible because of shared infrastructure, institutions, and social cooperation, and that a portion properly belongs to the collective that made the earning possible in the first place. Each can state the other's premise accurately, follow its internal logic, and recognize why someone starting from it would reach the conclusion they do. Neither is factually mistaken, and neither is reasoning in bad faith. What divides them is not a disputed fact about redistribution but a foundational disagreement about what individual earnings are, in the first place, owed to. This disagreement reflects how full understanding clarifies which foundational value is doing the work, without producing a shared value from which the dispute could be adjudicated.

Culinary Taste as Tier 2. Consider two people sharing the same two dishes at a restaurant: a delicately plated dish built around restrained seasoning and precise technique, and a rustic dish built around bold, assertive spicing and generous portions. One diner finds the delicate dish clearly superior, prizing balance and the way each flavor remains distinguishable rather than folded into a single dominant note; the other finds the rustic dish clearly superior, prizing intensity and the way bold flavors commit fully to one memorable impression rather than holding back. Each can explain their standard fully to the other, and neither disputes the plain facts of what is on the plate, only what good cooking is for, restraint and legibility on one side, commitment and intensity on the other. As with the painting, no shared culinary currency exists in

which precision and intensity could be weighed against each other to produce a single verdict, only two coherent, unranked standards of value applied to the same shared meal.

Culture Deep Dive

Culture is a useful case study precisely because it does not sit still in one tier.

Culture as Tier 1. In many cultures, declining a direct request outright is considered impolite, so a soft deferral such as “I will consider it” functions as a conventional way of saying no; someone accustomed to more direct communication may read the same phrase as genuine openness and press for a firmer answer later. This friction is Tier 1: once the convention is explained, both descriptions turn out to be correct, the phrase really was a polite refusal, and the listener really did hear an open sounding statement. Nothing either party values has to change, only what they know about how the other communicates.

Culture as Tier 3. Contrast this with a case that does not resolve. In some traditions, major personal decisions, such as who to marry or which career to pursue, are understood to properly rest with family or community authority rather than the individual alone. Someone reasoning from a tradition of individual autonomy can fully understand this framework, its internal logic, and the sense of belonging it provides, and still regard it as an unacceptable constraint on personal freedom. Someone reasoning from within that tradition can fully understand the autonomy based view and still regard it as a corrosive dissolving of the obligations that give a life its shape. Full understanding does not produce a shared picture the way it did with the indirect refusal; it clarifies that the two positions rest on different, comprehensive accounts of what a person owes to others, accounts that do not reduce to a shared covering standard. If tier were simply a function of subject matter, culture should behave consistently across cases of this kind. That it visibly does not, resolving in one instance and persisting in another, is itself evidence that a disagreement’s logical structure, not its topic, determines where it sits. The same conclusion follows for religion and politics considered on their own: not every religiously or politically inflected disagreement is Tier 3, but the ones that are cannot be handled the same way as the ones that are not, and subject matter alone will not tell you which is which.

Fashion as a Boundary Case. Clothing and personal style raise a question the case studies above leave open. Read as a matter of aesthetic taste alone, differing standards of dress, restrained and understated against bold and expressive, look like a **Tier 2** disagreement much like the painting or the meal, two coherent standards of value, no shared aesthetic currency, and no contradiction in holding both. But dress is rarely only aesthetic. Pressed further, a preference for modest, restrained clothing is, for some people, downstream of a broader conviction about self presentation, propriety, or religious obligation, in which case the disagreement inherits the structure of the conviction behind it rather than remaining a matter of style. Two people who each fully understand the other’s reasoning may still find that what looked like an unranked aesthetic difference is, once traced to its source, a **Tier 3** disagreement about what a person owes to modesty, community, or belief. This does not indicate an error in the taxonomy so much as a genuine feature of it, that the tier a disagreement occupies can depend on how far the reasoning behind a stated preference is traced, and that a single case, examined at different depths, can plausibly sit in more than one tier at once. Readers should take this as an invitation to apply the taxonomy carefully rather than mechanically, not as evidence that the three tiers collapse into one another.

	Tier 1: Resolvable-Perspectival	Tier 2: Interpersonally Incomparable	Tier 3: Doctrinally Incommensurable
Analogy	6 to you, 9 to me	One man's trash, another man's treasure	Agree to disagree
Structure	Same object, different vantage points	Same object, incomparable valuations	Different worldviews, no shared standard
Under full understanding	Resolves into one joint picture	Persists, not contradictory	Persists, cannot resolve
Formal anchor	none (null case)	Robbins (1932): interpersonal utility comparisons rest on no verifiable common unit	Rawls (1993): burdens of judgment produce reasonable, irreconcilable comprehensive doctrines
Register	Epistemic horizon	Structural	Ontological
Correctly scoped tool	AI mediated deliberation	Pareto preserving multi-objective RL	Constitutional AI, disclosed as non-neutral
Within tier failure	Aggregation	Scalarization	Presenting a declared choice as neutral
Implications	Shared context is sufficient, no representational workaround needed	Both valuations remain representable side by side, rather than ranked	No representation resolves the disagreement, only honest disclosure of one's own anchor remains available

Table 2: Detailed taxonomy of pluralistic disagreement across three tiers.