Date	Title	Quick Description	Link/Syllabus
Aug-2025	ENAIS AI Safety Atlas: Introduction to Technical AI Safety	An AI Safety introduction course which explores foundational safety concepts, such as alignment, scalable oversight, and evaluation strategies, to address the societal and existential challenges posed by superhuman intelligence.	https://ai-safety-atlas.com/
Aug-2025	CS6101: Rational Approaches to Cooperative Intelligence	A graduate-level research exploration module focused on scaling cooperative AI through rational modeling, Bayesian inference, and AI alignment.	https://cosilab.notion.site/cs6101-raci-fall-2025
Jan-2026	CS5340: Uncertainty Modelling in AI	A graduate-level course exploring probabilistic modeling methods for reasoning under uncertainty, with a primary focus on the representation, inference, and learning of Bayesian and Markov networks.	https://nusmods.com/courses/CS5340/uncertainty-modelling-in-ai

Schedule

Table

Week	Date	Topic	Readings
Week 1	13/08/2025	Introduction	Building Machines that Learn and Think With People; Open Problems in Cooperative AI
Week 2	20/08/2025	Theory of Mind	Bayesian Theory-of-Mind; The Naïve Utility Calculus
Week 3	27/08/2025	Inverse Planning	Planning-based Prediction for Pedestrians; Online Bayesian Goal Inference via Sequential Inverse Plan Search
Week 4	03/09/2025	Assistance	Cooperative Inverse Reinforcement Learning; AssistanceZero: Scalably Solving Assistance Games
Week 5	10/09/2025	Coordination and Teamwork	Exploring an Imagined "We" in Human Collective Hunting; Too Many Cooks: Bayesian Inference for Coordinating Multi-Agent Collaboration
Week 6	17/09/2025	Communication and Language	The Rational Speech Act Modeling Framework; Pragmatic Instruction Following and Goal Assistance via Cooperative Language-Guided Inverse Planning
Recess	24/09/2025	*RECESS*	
Week 7	01/10/2025	Teaching and Persuasion	A Rational Account of Pedagogical Reasoning; Bayesian Persuasion
Week 8	08/10/2025	Multi-Agent Games	Value-Function Reinforcement Learning in Markov Games; Stochastic Bayesian Games for Ad-Hoc Coordination
Week 9	15/10/2025	Norms	Social Norms as Choreography; Bayesian Norm Learning in Markov Games
Week 10	22/10/2025	Institutions	The Institutional Stance; What is Law? A Coordination Model of Legal Order
Week 11	29/10/2025	Negotiation and Deliberation	Automated Negotiation: Prospects, Methods and Challenges; The Habermas Machine
Week 12	05/11/2025	Alignment	Beyond Preferences in AI Alignment; Beyond Alignment: The Patchwork Quilt of Human Coexistence
Week 13	12/11/2025	Project Presentations	

Week	Date	Lecture
1	13 Jan 2026	Probability Basics
2	20 Jan 2026	Simple Probabilistic Models
3	27 Jan 2026	Bayesian Networks (directed)
4	3 Feb 2026	Markov Random Fields (undirected)
5	10 Feb 2026	Inference on Networks
6	17 Feb 2026	Factor Graphs
Recess	24 Feb 2026	N/A
7	3 Mar 2026	Latent Variables and the EM algorithm
8	10 Mar 2026	Hidden Markov Model
9	17 Mar 2026	Approximate Inference via Sampling, MCMC
10	24 Mar 2026	Variational Inference
11	31 Mar 2026	Inference and Decision-Making
12	7 Apr 2026	Closing Lecture
13	14 Apr 2026	Presentations

Apr-2026	Technical Alignment Research Accelerator (TARA)	A 14-week intensive program which utilizes the ARENA curriculum to equip researchers and engineers with technical AI safety expertise through the practical implementation of transformers, mechanistic interpretability, and reinforcement learning.	https://www.taraprogram.org/curriculum	**Week** — **Material Covered** Week 0 - March 7 — Prerequisites + PyTorch practice Week 1 - March 14 — CNNs and ResNets Week 2 - March 21 — Optimisation Week 3 - March 28 — Transformer from Scratch Week 4 - April 4 — Intro to Mechanistic Interpretability Week 5 - April 11 — Linear Probes Week 6 - April 18 — Intro to RL Week 7 - April 25 — Proximal Policy Optimisation Week 8 - May 2 — RLHF Week 9 - May 9 — Intro to Evals + Dataset Generation Week 10 - May 16 — Running Evals with Inspect Week 11 - May 23 — LLM Agents Week 12 - May 30 — Project Week I Week 13 - June 6 — Project Week II Week 14 - June 13 — Project Presentations		
May-2026	BlueDot Technical AI Safety	A technical course providing an in-depth exploration of the engineering frameworks required to build safe AI, focusing on model alignment, safety evaluations, and mechanistic interpretability to mitigate the risks associated with frontier models.	https://bluedot.org/courses/technical-ai-safety			